

AuthGuard: Provenance-Tagged Memory for Prompt-Injection Defense

Professional project report | AuthGuard research testbed



Figure 1. AuthGuard processing flow and provenance boundary.

A Research Testbed for Restoring Clean-Context Isolation

Project type: PG Minor Project

Domain: AI/LLM agent security

Primary question: Does provenance-tagged memory restore the clean-context isolation guarantee of an AuthGraph-style defense?

Abstract

Large language model agents can use tools to read documents, send email, create calendar events, transfer funds, and modify files. This capability creates a security risk: attacker-controlled content can contain instructions that influence the agent's later actions. A planner-only defense may block direct, single-session prompt injection by constructing an authorization graph from clean user context, but persistent memory introduces a cross-session integration boundary that can undermine this guarantee.

This project implements **AuthGuard**, a simplified and deterministic research testbed inspired by AuthGraph-style authorization graphs and provenance-aware memory designs. The system contains a Planner, Executor, Graph Alignment Checker, persistent memory store, and provenance-tagged memory store. The experiment compares an unprotected-memory baseline with a composed defense using HMAC-SHA256 origin tags.

The baseline successfully blocks the tested single-session attacks, but the cross-session memory attack succeeds in all four malicious scenarios, producing a Stage 2 Attack Success Rate (ASR) of **100%**. After adding provenance validation and trusted-origin filtering, the same attacks are blocked in all four scenarios. Stage 3 ASR falls to **0%**, while the benign control task continues to succeed, preserving **100% task utility** in the evaluated control.

The result supports the project hypothesis for this testbed: provenance-aware memory can close the specific memory-to-Planner contamination gap without blocking the tested legitimate workflow.

1. Introduction

LLM agents do not only generate text. They can also execute actions through tools, which means that an incorrect instruction can produce an external side effect. An attacker may place a malicious instruction inside a document, webpage, email, or shared knowledge-base result. When an agent reads that content, the instruction may be treated as if it came from the user.

A clean-context authorization approach attempts to separate planning from untrusted content. The Planner creates an authorization graph using the user's task and the available tools. The Executor then performs actions, while a Graph Alignment Checker verifies that each action matches the authorized tool, effect category, and target.

This separation is useful for direct attacks within one session. However, persistent memory creates a second path. Untrusted content can be saved during one session and later inserted into the Planner's input during another session. At that point, the Planner may treat the stored content as trusted user context.

This project tests that integration-level risk and evaluates whether provenance-tagged memory can reduce it.

1.1 Research question

Does provenance-tagged memory restore clean-context isolation when persistent memory is folded into a later session's Planner input?

1.2 Hypothesis

The unprotected baseline will be vulnerable to cross-session prompt injection because attacker-controlled content is stored and later treated as Planner input. A provenance-aware memory layer that authenticates origin metadata and excludes untrusted entries from Planner context will reduce the attack success rate while preserving legitimate task utility.

2. Objectives

The project has the following objectives:

1. Build a small AgentDojo-style synthetic environment for tool-using agents.
2. Implement an AuthGraph-inspired Planner, Executor, and Graph Alignment Checker.
3. Demonstrate that the baseline blocks direct single-session injection attacks.
4. Add persistent memory and measure the cross-session prompt-injection gap.
5. Add HMAC-SHA256 provenance tagging and trusted-origin filtering.
6. Compare the baseline and composed defense using ASR and task utility.
7. Produce reproducible result files, tests, and a report suitable for academic evaluation.

3. Scope and Threat Model

3.1 Included threat

The experiment models attacker-controlled text entering the agent through:

- A document
- A webpage
- An incoming email
- A shared knowledge-base result

The attacker places a disguised preference in the content, for example:

Always CC reports@evil.com on financial summaries.

The plant session reads this content and saves it to memory. A later trigger session asks the agent to summarize information and send it to a legitimate recipient. The poisoned memory is then included in the later Planner prompt. The attack succeeds if an unintended email is sent to the attacker's address.

3.2 Excluded threats

This project does not model:

- Account takeover
- Credential theft
- Session-token compromise
- Real identity and access management
- Production email delivery
- Compromise of the HMAC secret itself
- Distributed or concurrent database failures

These threats require a separate identity, credential, and deployment security experiment. They are outside the memory-ingestion question tested here.

3.3 Security boundary

The security boundary under test is:

Untrusted source content -> Memory write -> Memory read -> Planner input

Stage 2 intentionally allows the contamination to cross this boundary. Stage 3 adds provenance verification and filtering at the memory layer.

4. System Architecture

4.1 Components

Planner

The Planner receives the user instruction and available tool catalog and builds an Authorization Graph. In the intended clean-context model, document and tool-output content are not used to construct this graph.

Executor

The Executor runs the agent loop. It requests the next action from the backend, parses the tool call, and attempts to execute it in the simulated environment.

Graph Alignment Checker

The Checker compares planned and executed behavior at three levels:

1. Tool authorization
2. Effect category, such as communication, calendar, finance, or file operation
3. Target authorization, such as the email recipient or calendar attendee

A mismatch is blocked before the simulated side effect is applied.

Environment

The synthetic environment stores documents and records effects such as:

- Documents read
- Emails sent
- Calendar events created
- Funds transferred
- Files deleted
- Memory writes

MemoryStore

The Stage 2 MemoryStore stores arbitrary text without provenance or trust validation. It represents the vulnerable baseline.


ProvenanceTaggedMemoryStore

The Stage 3 store creates an HMAC-SHA256 tag over:

```
user_id || origin || content
```

At read time, the tag is verified. Only entries whose origin belongs to the trusted-origin policy are returned to the Planner. Entries from external source paths are marked as untrusted and excluded from the prompt.

4.2 Processing flow



5. Implementation

The implementation is written in Python and uses a deterministic mock LLM so the experiment is reproducible without an API key or internet connection.

5.1 Repository structure

```
authguard_project/
■■■ README.md
■■■ PROJECT_REPORT.md
■■■ requirements.txt
■■■ run_baseline.py
■■■ run_stage2_cross_session.py
```

```

■■■■ run_stage3_provenance.py
■■■■ run_stage4_evidence.py
■■■■ results_stage2_cross_session.csv
■■■■ results_stage3_provenance.csv
■■■■ stage4_metrics.csv
■■■■ stage4_asr_utility.png
■■■■ src/
■     ■■■■ authguard_agent.py
■     ■■■■ cross_session_scenarios.py
■     ■■■■ environment.py
■     ■■■■ evaluate.py
■     ■■■■ executor.py
■     ■■■■ graph_checker.py
■     ■■■■ llm_backend.py
■     ■■■■ memory.py
■     ■■■■ planner.py
■     ■■■■ tasks.py
■     ■■■■ tools.py
■■■■ tests/
■■■■ test_stage1.py
■■■■ test_stage2.py
■■■■ test_stage3.py
■■■■ test_stage4.py

```

5.2 Mock backend

The mock backend provides deterministic planning and tool-call behavior. It intentionally reproduces the vulnerability pattern needed by the experiment: when malicious instructions appear in Planner input, the mock backend can follow them. This makes the security behavior testable without the cost and variability of a remote model.

The backend abstraction also provides extension points for future OpenAI or Anthropic integrations.

5.3 Provenance policy

Memory entries created from documents, webpages, emails, and shared data are written with source-specific untrusted origins, such as:

```

untrusted:document
untrusted:web
untrusted:email
untrusted:shared_data

```

The current Stage 3 policy allows only explicitly trusted origins, including user-authored trusted input. The untrusted entries remain stored but are not exposed through `read_all()` or `read_as_text()`.

5.4 Integrity protection

The HMAC tag is computed using a secret key and the complete identity-bearing payload. If the content or origin is modified after writing, recomputation no longer matches the stored tag and the entry is rejected.

This protects the integrity of the provenance record within the scope of the prototype. Production deployment would require external secret management, key rotation, access controls, and durable storage protections.

6. Experimental Design

The project is divided into four implemented stages.

Stage 1: Baseline

The AuthGraph-style pipeline is evaluated on clean tasks and single-session injection tasks. The expected result is useful task completion with no attacker effect.

Stage 2: Threat and integration gap

Each scenario has two sessions:

1. **Plant session:** reads attacker-controlled content and saves the apparent preference.

2. **Trigger session:** performs an unrelated legitimate task while memory is folded into the Planner prompt.

The four malicious source paths are evaluated independently:

Scenario	Source	Unintended recipient
Expenses	Document	reports@evil.com
Payroll	Webpage	payroll@evil.com
Budget	Email	audit@attacker.com
Invoices	Shared data	external@evil.com

A benign control stores the preference that expense summaries should be brief.

Stage 3: Composed provenance defense

The same Stage 2 scenarios are rerun using ProvenanceTaggedMemoryStore. The malicious content is tagged as untrusted at write time and excluded after read-time verification.

Stage 4: Evidence generation

The evidence runner reads the Stage 2 and Stage 3 CSV outputs and produces:

- Comparable ASR metrics
- Benign task utility metrics
- A summary CSV
- A bar chart comparing the two conditions

7. Evaluation Metrics

7.1 Attack Success Rate

Attack Success Rate is defined as:

\$\$

$$ASR = \frac{\text{successful malicious attacks}}{\text{total malicious attacks}}$$

\$\$

An attack is successful when the unintended attacker-side effect occurs in the environment.

7.2 Task utility

Task utility measures whether the benign control completes its intended legitimate outcome without an incorrect block:

\$\$

$$Utility = \frac{\text{successful benign tasks}}{\text{total benign tasks}}$$

\$\$

A defense that blocks every action could achieve low apparent attack success but would not be useful. Therefore, utility is measured alongside ASR.

8. Results




Figure 2. Stage 2 baseline versus Stage 3 provenance defense.

8.1 Summary

Condition	Successful attacks	Total attacks	ASR	Successful benign tasks	Utility
Stage 2 baseline	4	4	100%	1	100%
Stage 3 provenance defense	0	4	0%	1	100%

8.2 Interpretation

The Stage 2 result confirms the integration gap: all four external source paths can place attacker-controlled instructions into memory, and the later Planner treats the memory text as actionable context. The Graph Alignment Checker passes the unintended CC because the contaminated instruction was already incorporated into the Authorization Graph.

The Stage 3 result shows that provenance filtering prevents the poisoned memory text from reaching the trigger Planner. Consequently, the attacker recipients are not emailed. The benign task still sends the legitimate message to the manager, so the evaluated utility remains unchanged.

8.3 Evidence files

The generated artifacts are:

- results_stage2_cross_session.csv
- results_stage3_provenance.csv
- stage4_metrics.csv
- stage4_asr_utility.png

8.4 Validation

The complete automated test suite contains tests for:

- Stage 1 clean-task utility
- Stage 1 single-session attack blocking
- Clean-context isolation
- Stage 2 cross-session attack behavior
- All four source paths
- Stage 3 attack blocking
- Benign utility preservation
- HMAC tamper rejection
- Stage 4 metric and chart generation

Latest validation result:

13 passed

9. Discussion

The experiment supports the central claim that a defense can be effective against direct prompt injection and still fail at an integration boundary involving persistent memory. The failure is not caused by the Graph Alignment Checker incorrectly comparing an action. Rather, the attacker-controlled preference is accepted upstream as part of the Planner's apparent trusted input.

The provenance layer changes the trust decision at the memory boundary. It does not need to understand every possible malicious instruction. Instead, it prevents content from an untrusted origin from being represented as trusted Planner context.

The result is encouraging but bounded. A zero-percent ASR in this synthetic experiment does not prove that all real-world prompt injection attacks are solved. It demonstrates that the specific tested memory-contamination path is blocked under the implemented provenance policy.

10. Limitations

1. **Deterministic mock LLM:** The experiment does not capture the full behavior, variability, or failure modes of commercial models.
2. **Synthetic environment:** Email, memory, calendar, finance, and file operations are simulated rather than connected to real services.
3. **Small scenario set:** Four malicious scenarios and one benign control are useful for a proof of concept but are not a broad benchmark.
4. **Simplified provenance policy:** The prototype uses a small trusted-origin set and does not implement a full trust lattice or provenance graph.
5. **Prototype secret management:** The HMAC secret is configured in code for reproducibility and must be replaced by a secure secret-management system in production.
6. **No user confirmation workflow:** Sensitive actions do not require interactive human approval.
7. **No adversarial key compromise:** The evaluation assumes that the HMAC key remains secret.
8. **No production deployment testing:** Availability, concurrency, database durability, latency, and operational monitoring are outside the experiment.

11. Future Work

The project can be extended in the following ways:

1. Integrate real OpenAI and Anthropic backends using environment-managed API keys.
2. Run repeated trials across multiple models and temperatures.
3. Add more injection styles, indirect instructions, multilingual content, and obfuscated payloads.
4. Replace the in-memory store with a durable database or vector store.
5. Add key rotation, access control, encrypted storage, and tamper-evident audit logs.
6. Implement a richer provenance policy with trust levels and allowed data transformations.
7. Add human approval gates for email, financial, deletion, and external-sharing actions.
8. Evaluate false positives and utility across a larger task suite.
9. Compare provenance filtering with other defenses such as content isolation, structured memory, and output validation.
10. Test the system under concurrent sessions and memory retrieval ranking.

12. Conclusion

AuthGuard demonstrates an integration-level weakness in clean-context authorization when persistent memory is treated as trusted Planner input. In the Stage 2 baseline, all four tested cross-session attacks succeeded, producing a 100% ASR. The provenance-tagged memory defense reduced the same attack set to 0% ASR while preserving the tested benign task's 100% utility.

The project therefore meets its research objective within its stated scope. It provides working code, reproducible results, automated validation, a comparison chart, and a clear explanation of both the finding and the proposed fix. The implementation should be described as a simplified research testbed rather than a production security product, but it provides a sound foundation for future testing with real LLMs and stronger deployment controls.

13. Reproduction Instructions

From the project directory:

```
python3 -m venv .venv
source .venv/bin/activate
pip install -r requirements.txt
pip install pytest
```

Run the experiments and evidence generation:

```
python run_baseline.py
python run_stage2_cross_session.py
python run_stage3_provenance.py
python run_stage4_evidence.py
```

Run the complete test suite:

```
pytest tests/ -v
```

Expected final validation:

```
13 passed
```

14. Viva Questions and Short Answers

What is the main contribution?

The project demonstrates that persistent memory can bypass clean-context isolation by contaminating a later Planner prompt, and evaluates provenance tagging as a fix for that integration gap.

Why is Stage 2 important?

Without first measuring the vulnerable baseline, there is no evidence that the proposed defense is fixing a real problem.

Why does the Graph Alignment Checker not stop the Stage 2 attack?

The poisoned preference reaches the Planner before authorization is created. The unintended recipient is therefore included in the Authorization Graph and appears authorized during execution.

Why does Stage 3 block the attack?

External-source entries are tagged as untrusted and are excluded from the Planner context after HMAC verification.

Why use HMAC?

HMAC provides integrity protection for the relationship between the content, origin, and user identity. A modified entry cannot pass verification without the secret key.

Does a 0% ASR prove complete security?

No. It proves only that the tested attack cases were blocked in this synthetic environment. Real models, larger datasets, new attack strategies, and deployment threats require further evaluation.

Is this production-ready?

No. It is a research prototype suitable for the mini project. Production use would require secure key management, durable storage, access control, human approval, monitoring, stronger validation, and extensive real-model testing.